

ENME 808 – HW 1

Adrienne Rudolph

02/16/2024

Question 1

```
# Author: Adrienne Rudolph
# Class: ENME 808
# HW1 Problem 1

import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
from sklearn.linear_model import LinearRegression
from sklearn.metrics import root_mean_squared_error as rootMSE

#Read in the data
mlr05 = pd.read_excel("mlr05.xlsx")

#Prepare the data
X_varb = mlr05.iloc[:, 1:] #X variables are the X2, X3, .. X6 predictors
Y_varb = mlr05.iloc[:, 0]  #Y variable is the sales data, X1

#Separate Training and Testing Sets
x_train = X_varb[:20] #training set are X_varb columns, first 20 rows
y_train = Y_varb[:20] #training set is 'sales data' column, first 20 rows

x_test = X_varb[20:] #test set are all columns, last 7 rows
y_test = Y_varb[20:] #test set is 'sales data' column, last 7 rows

#Select the linear regression model
dataModel = LinearRegression()

#Train the model
dataModel.fit(x_train, y_train)

#Make a prediction for the sales data
prediction = dataModel.predict(x_test)


#Evaluate for error in actual and prediction
rmse = rootMSE(y_test, prediction)
print(rmse)
```

```

#Reshape and print prediction and
prediction = prediction.reshape(-1,1) #Visually appealing column vector for
printing purpose
print(prediction)

#Plot test and prediction, and show line of perfect prediction
plt.scatter(y_test, prediction, label='Predicted vs Actual')
plt.xlabel('Actual Sales Data')
plt.ylabel("Predictied Sales Data")
plt.title('Actual vs Predicted Sales Data Values')
plt.plot([min(y_test), max(y_test)], [min(y_test), max(y_test)], color='purple',
linestyle='--')
plt.legend()
plt.grid(True)
plt.show()

```



The root mean squared error between the actual and predicted data is 59.8, which doesn't seem to be a very good outcome. This means the predictor doesn't model the data very well, and indicates the data likely requires a nonlinear predictor. The purple dotted line shown in the plot above displays a perfect match between predicted and actual sales data.

Adrienne Rudolph

HW 1

Question 2

a.)

Perceptron 2 dimensions

$$h(x) = \text{sign}(\omega^T x)$$

$$\omega = [\omega_0, \omega_1, \omega_2]^T$$

$$x = [1, x_1, x_2]^T$$

$$h(x) = +1 \text{ when } \omega^T x > 0$$

$$h(x) = -1 \text{ when } \omega^T x < 0$$

$$\text{boundary at } \omega^T x = 0$$

$$\text{so } \omega_0 + \omega_1 x_1 + \omega_2 x_2 = 0$$

$$\omega_1 x_1 + \omega_2 x_2 = -\omega_0$$

$$\omega_2 x_2 = -\omega_0 - \omega_1 x_1$$

$$x_2 = -\frac{\omega_0}{\omega_2} - \frac{\omega_1}{\omega_2} x_1$$

$$x_2 = -\frac{\omega_1}{\omega_2} x_1 - \frac{\omega_0}{\omega_2}$$


$$\text{where } a = -\frac{\omega_1}{\omega_2} \quad \text{and } b = -\frac{\omega_0}{\omega_2}$$

b.) $\omega = [1, 2, 3]^T$ $\omega_0 = 1, \omega_1 = 2, \omega_2 = 3$
 $x_2 = -\frac{2}{3}x_1 - \frac{1}{3}$


$\omega = [-1, -2, -3]^T$ $\omega_0 = -1, \omega_1 = -2, \omega_2 = -3$
 $x_2 = -\left(-\frac{2}{3}\right)x_1 - \left(-\frac{1}{3}\right)$
 $x_2 = -\frac{2}{3}x_1 - \frac{1}{3}$

Both cases have the same line, but the regions change

case 1: $[1, 2, 3]^T$



case 2: $-[1, 2, 3]^T$



Question 3

- linearly separable dataset
- $B = \text{minimum} \{ \|\omega\| : \forall i \in [m], y_i (\omega^T x_i) \geq 1 \}$
- $R = \max_i \|x_i\|$

$$\omega^{t+1} = \omega^t + y_i x_i$$

optimal ω^* must be equal to or greater than the $\min \{ \}$, or $\geq B$.

** $\omega(t+1)$ will always be closer to ω^* than $\omega(t)$*
assume $\|\omega^*\| = 1$

$$\begin{aligned} \omega^{t+1} \omega^* &= (\omega^t + (y_i x_i)) \omega^* \\ &= \omega^t \omega^* + y_i (\omega^* x_i) \rightarrow B \\ &> \omega^t \omega^* + B \quad \leftarrow \text{Because minimum} \end{aligned}$$

so after a 't' number of iterations
 $\omega^{t+1} \cdot \omega^* > tB$

$$\begin{aligned} \text{Because } \omega^{t+1} \cdot \omega^* &\leq \|\omega^{t+1}\| \|\omega^*\| = \|\omega^{t+1}\| \\ \|\omega^{t+1}\| &> tB \quad \text{lower bound} \end{aligned}$$

To get an upper bound

$$\begin{aligned} \rightarrow \|\omega^{t+1}\|^2 &= \|\omega^t + y_i x_i\|^2 \\ &= \|\omega^t\|^2 + 2\|y_i x_i\|^2 + 2y_i (\omega^t \cdot x_i) \\ &= \|\omega^t\|^2 + \|x_i\|^2 + \underbrace{2y_i (\omega^t \cdot x_i)}_{\substack{\text{goes away for} \\ \text{upper bound}}} \\ &\leq \|\omega^t\|^2 + \|x_i\|^2 \end{aligned}$$

$$\|x_i\|^2 = R^2$$

$$\rightarrow \leq \|\omega^t\|^2 + R^2$$

so after t iterations: $\|\omega^{t+1}\|^2 \leq tR^2$

$$t^2 B^2 < \|\omega^{t+1}\|^2 \leq tR^2$$

upper bound

$t \leq \frac{R^2}{B^2}$ stops after at most

Question 4

```
import numpy as np
import matplotlib.pyplot as plt

#pick a function x(t), y(t)
def targetFunc(x):
    return 3*x + (1/2)

#Produce the randomly generated data
dataPts = np.random.uniform(-10, 10, size=(10100, 2))

#Separate Train and Test Data
trainPts = dataPts[:100, :]
testPts = dataPts[100:, :]

#Separate X and Y columns from Train Data
train_feature = trainPts[:, 0] #x data for training set
train_label = trainPts[:, 1]   #y data for training set

#Separate X and Y columns from Test Data
test_feature = testPts[:, 0]   #x data for test set
test_label = testPts[:, 1]     #y data for test set

#Initialize w vector and eta (n)
w = np.zeros((3, 1))
n = 0.0001
y_train_label = []

#Iterate, calculate signal, and update weights
for i in range(0, 10):
    for t in range(len(trainPts)):
        y_actual = trainPts[t,1] - targetFunc(trainPts[t,0])
        if y_actual > 0:
            y_actual = 1
        else:
            y_actual = -1
        if i == 0:
            y_train_label.append(y_actual)

        x_vec = np.array([1, trainPts[t,0], trainPts[t,1]]).reshape(-1,1)

        s_t = np.sign(np.dot(w.T, x_vec))

        temp = y_actual * s_t
```

```

        if temp <= 1:
            w += n * (y_actual - s_t) * x_vec
w = w

w0 = w[0]
w1 = w[1]
w2 = w[2]
def output(z):
    return (-(w1/w2) * z) - (w0/w2)

#Apply best weights to test set
predictions = []
actual = []
for a in range(len(testPts)):
    y_actual_test = testPts[a,1] - targetFunc(testPts[a,0])
    if y_actual_test > 0:
        y_actual_test = 1
    else:
        y_actual_test = -1

    actual.append(y_actual_test)
    x_vec = np.array([1, test_feature[a], test_label[a]]).reshape(-1,1)

    s_t = np.sign(np.dot(w.T, x_vec))
    predictions.append(s_t)

y_test_label = actual
predictions = np.array(predictions).flatten()
actual = np.array(actual).flatten()
print(np.sum(predictions==actual)/predictions.shape[0])

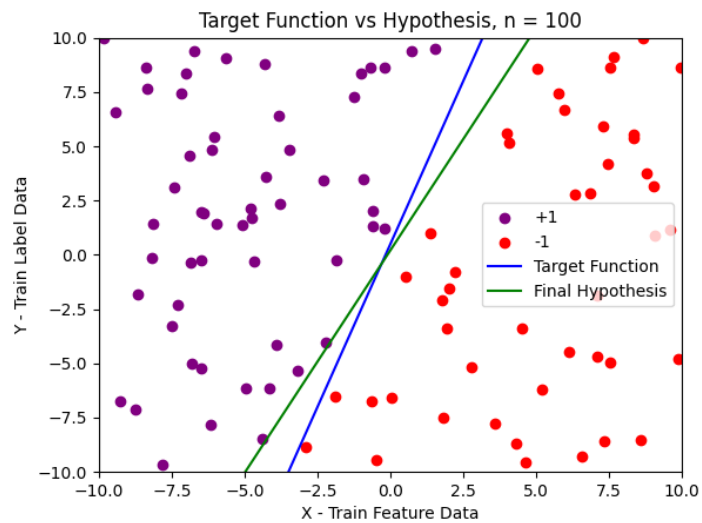
y_train_label = np.array(y_train_label).flatten().astype(int)
y_test_label = np.array(y_test_label).flatten().astype(int)

# Plot the training data set
z = np.linspace(-10,10)


plt.scatter(train_feature[y_train_label==1], train_label[y_train_label==1],
color='purple',label='+1')
plt.scatter(train_feature[y_train_label==-1], train_label[y_train_label==-1],color='red',label='-1')
plt.plot(z, targetFunc(z), color='blue', label='Target Function')
plt.plot(z, output(z), color='green', label='Final Hypothesis')
plt.title('Target Function vs Hypothesis, n = 0.0001')

```

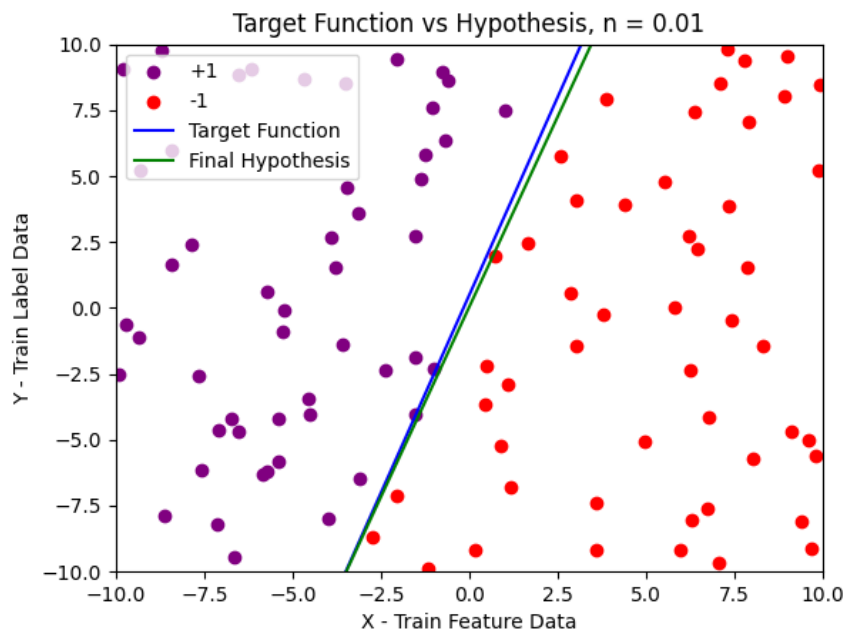
```
plt.xlabel('X - Train Feature Data')
plt.ylabel('Y - Train Label Data')
plt.legend()
plt.xlim(-10,10)
plt.ylim(-10,10)
plt.show()
```




The accuracy of the hypothesis on the test set in the case of $\eta = 100$ is 95.97%.



The accuracy of the hypothesis on the test set in the case of $\eta = 1$ is 99.58%



The accuracy of the hypothesis on the test set in the case of $\eta = 0.01$ is 99.4%



The accuracy of the hypothesis on the test set in the case of $\eta = 0.0001$ is 98.96%

It is somewhat difficult to comment on just how different the results are based on the learning rate, η . The data set I generated is random, so each time I run the program with a new learning rate, the overall data set and accuracy will change a little bit. However, in general, it does appear that a smaller learning rate might help the algorithm produce a more accurate hypothesis.

Question 5

Exercise 1.8 - If $\mu = 0.9$, what is the probability that a sample of 10 marbles will have $V \leq 0.1$?

Hint: use binomial distribution

The answer is a very small number

$$\binom{n}{K} \mu^K (1-\mu)^{n-K}$$

$$\Pr(V \leq 0.1) \quad \text{where } V = K/n$$

$$\rightarrow \Pr(V = \frac{K}{n} \leq 0.1) \rightarrow \Pr(K \leq 0.1 \cdot n)$$

$$\rightarrow \Pr(K \leq 1)$$

$$\text{so } \Pr(0) + \Pr(1) = \sum_0^1 \binom{n}{x} p^x (1-p)^{n-x}$$

$$\text{where } \binom{n}{x} = \frac{n!}{x!(n-x)!}$$

$$= \binom{10}{0} \mu^0 (1-\mu)^{10-0} + \binom{10}{1} \mu^1 (1-\mu)^{10-1}$$

$$= \frac{10!}{0!(10-0)!} (0.9)^0 (1-0.9)^{10-0} + \frac{10!}{1!(10-1)!} (0.9)^1 (1-0.9)^{10-1}$$

$$= 1 \times 10^{-10} + 9 \times 10^{-9}$$

$$\Pr(V \leq 0.1) = 9.1 \times 10^{-9}$$

Exercise 1.9

$$\mu = 0.9$$

sample of 10 marbles

$$\nu \leq 0.1$$

$$P[|\nu - \mu| > \epsilon] \leq 2e^{-2\epsilon^2 N}$$

$$P[|0.1 - 0.9| > \epsilon] \leq 2e^{-2\epsilon^2 N}$$

$$P[0.8 > \epsilon] \leq 2e^{-2\epsilon^2 N}$$

so plug in 0.8 for ϵ

$$P[\nu \leq 0.1] \leq 2e^{-2(0.8)^2(10)}$$

$$P[\nu \leq 0.1] \leq 5.522 \times 10^{-6}$$

The probability of the "bad" event happening is higher coming from the Hoeffding Inequality than the binomial distribution. That means this calculation is more conservative.